

Produced by:

Dr. Brenda Mullally

bmullally@wit.ie

Department Computing and Maths

Waterford Institute of Technology

www.wit.ie

moodle.wit.ie

HDip Business Systems Analysis Data Analytics

Introduction

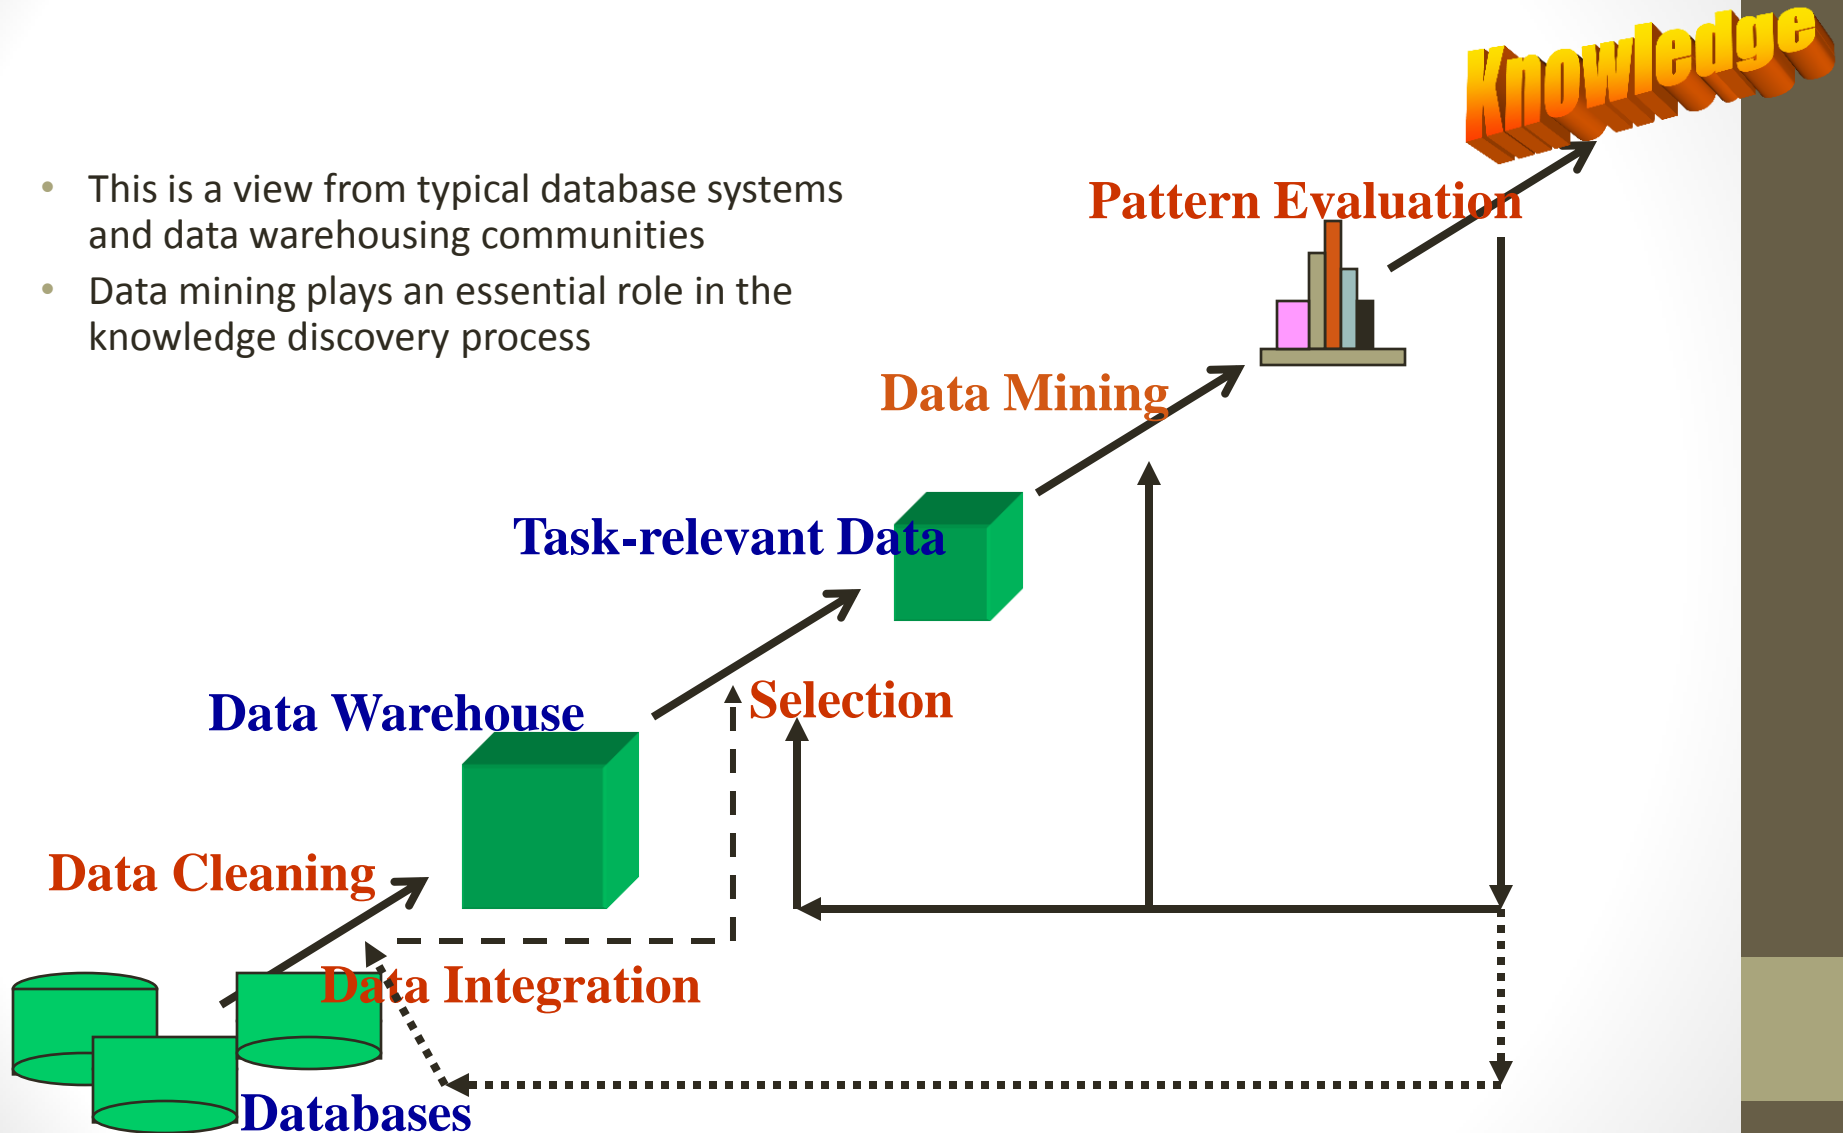
- Business analytics
- Data mining
- Goal, design, techniques & implementation of data mining (CRISP-DM)
- Data mining application
- Understand data

Business Analytics

- The essence of business analytics is to embed analytical decisions into business processes on an ongoing, repeatable basis instead of treating analytics as an ad hoc activity (Davenport and Harris, 2007).
- Relevant technologies for business analytics, including different data mining technologies that have their origin in statistics and artificial intelligence are covered in this section of the module.
- Firms need to follow a business analytics strategy to embed analytics in their most important business processes.

Knowledge Discovery from Data (KDD) Process

- This is a view from typical database systems and data warehousing communities
- Data mining plays an essential role in the knowledge discovery process



Data Mining

- IVF - 60 features
- Dairy farmers – 700 attributes
- Plant diseases – 35 attributes
- Life and death.

Data Mining

- Overwhelmed with data.
- Ubiquitous computing
- Swipe cards
- WWW
- Volume increasing – proportion understand reduced
- Patterns?
- Profit?

Technologies to Create Insights: Using Data Mining to Create New Explicit Knowledge

- Business analytics
- *Knowledge discovery in databases* (KDD), or more commonly, *data mining* (DM)
- Analytics comes with hypothesis testing
- Data mining is more the act of discovery that lacks a hypothesis
- DM applications have made noteworthy contributions to scientific discovery, for example, in breast cancer diagnosis

Data Mining

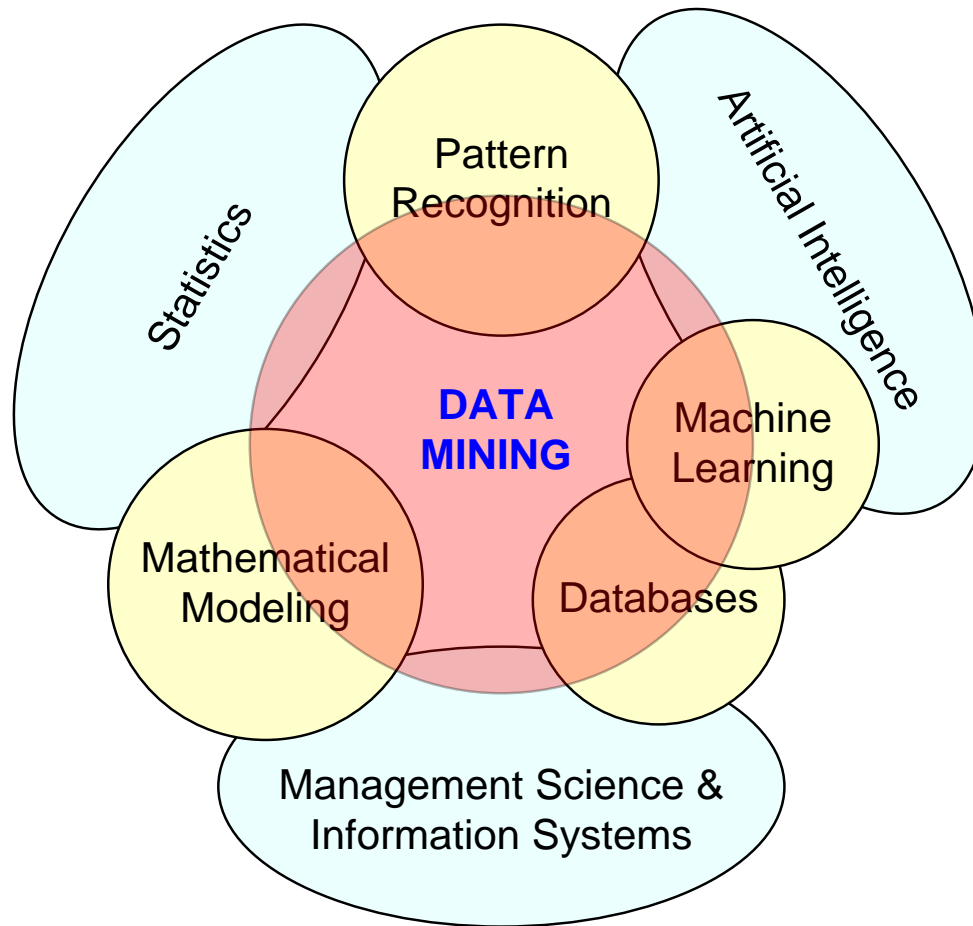
- The nontrivial process of identifying valid, novel, potentially useful, and ultimately understandable patterns in data stored in structured databases - Fayyad et al., (1996)
- Keywords in this definition: Process, nontrivial, valid, novel, potentially useful, understandable
- Data mining: a misnomer?
- Other names: knowledge extraction, pattern analysis, knowledge discovery, information harvesting, pattern searching, data dredging

DM characteristics/objectives

- Data often buried deep in large databases with several years of data.
- Data in a variety of formats
- Usually client server architecture
- Stored electronically
- Automated/semi-automated search
- Used for prediction
- Analysing data already present in databases.
- Meaningful patterns - economic advantage
- Structural patterns
- Machine learning

Data Mining at the Intersection of Many Disciplines

A process that uses statistical, mathematical, artificial intelligence and machine-learning techniques to extract and identify useful information and subsequent knowledge from large databases



Data mining commercial view

- Lots of data is being collected:
 - E-commerce, web data
 - Any time you go to a grocery store
 - Any time you use your bank or credit card
- Computers are cheaper and more powerful
- Competitive pressure is strong

Data Mining Applications

- Travel industry
- Healthcare
- Medicine
- Entertainment industry
- Homeland security
- Sports
- Crime

Data Mining Applications

- Brokerage and Securities Trading
 - Predict changes on certain bond prices
 - Forecast the direction of stock fluctuations
 - Assess the effect of events on market movements
 - Identify and prevent fraudulent activities in trading
- Insurance
 - Forecast claim costs for better business planning
 - Determine optimal rate plans
 - Optimize marketing to specific customers
 - Identify and prevent fraudulent claim activities

Data Mining Applications

- Customer Relationship Management
 - Maximize return on marketing campaigns
 - Improve customer retention (churn analysis)
 - Maximize customer value (cross- or up-selling)
 - Identify and treat most valued customers
- Banking & Other Financial
 - Automate the loan application process
 - Detecting fraudulent transactions
 - Maximize customer value (cross- and up-selling)
 - Optimizing cash reserves with forecasting

Data Mining Applications

- Retailing and Logistics
 - Optimize inventory levels at different locations
 - Improve the store layout and sales promotions
 - Optimize logistics by predicting seasonal effects
 - Minimize losses due to limited shelf life
- Manufacturing and Maintenance
 - Predict/prevent machinery failures
 - Identify anomalies in production systems to optimize manufacturing capacity
 - Discover novel patterns to improve product quality


Data Mining Process

- A manifestation of best practices
- A systematic way to conduct DM projects
- Most common standard processes:
 - CRISP-DM (Cross-Industry Standard Process for Data Mining)
 - SEMMA (Sample, Explore, Modify, Model, and Assess)
 - KDD (Knowledge Discovery in Databases)

Data Mining Process

What main methodology are you using for your analytics, data mining, or data science projects ? [200 votes total]

2014 poll 2007 poll



Source: KDNuggets.com, Oct 2014

Data Mining Process: CRISP-DM



Data Mining Process: CRISP-DM

Step 1: Business Understanding

Step 2: Data Understanding

Step 3: Data Preparation (!)


Step 4: Model Building

Step 5: Testing and Evaluation

Step 6: Deployment

- The process is highly repetitive and experimental

Data Preparation – A Critical DM Task



Output:

Dataset and a dataset description, both used for modeling or for the major analysis work of the project.

Data Preparation – A critical task

Select Data

- Activities:
 - Source data collected.
 - Significance and correlation tests to determine inclusion
 - Data selection criteria revisited (step2)
 - Data subsets
 - Sampling techniques
 - Document rationale for inclusion/exclusion
- Output:
 - List of data to be used/excluded and reasons

Data Preparation – A critical task

Clean Data

- Activities:
 - Noise
 - Special values.
 - Data selection criteria
- Output:
 - Data cleaning report

Data Preparation – A critical task

Construct Data

- Activities:
 - Construction mechanisms and tools
 - Production of derived attributes, new records, transformed values
 - Selection criteria
- Output:
 - Derived Attributes

Data Preparation – A critical task

Integrate Data

- Activities:
 - Integrate sources
 - Data selection criteria
- Output:
 - Merged data

Data Preparation – A critical task

Format Data

- Activities:
 - Rearranging attributes
 - Reordering records
 - Reformatting values
- Output:
 - Reformatted data